

COMP2501 Assignment 2

Sibo Ding

Spring 2023

Requirements

Submission deadline: Mar 28th, 2023 at 23:59.

Full mark of assignment 2: 33.

For the following questions, please:

1. Replace all [Input here] places with your information or your answer.
2. Complete the code block by adding your own code to fulfill the requirements in each question. Please use the existing code block and do not add your own code block. Noting that please use `head()` to show the corresponding results if there are too many rows in them.

Please make sure your Rmd file is a valid Markdown document and can be successfully knitted.

For assignment submission, please knit your final Rmd file into a Word document, and submit both your **Rmd** file and the knitted **Microsoft Word** document file to Moodle. You get 0 score if 1) the Rmd file you submitted cannot be knitted, and 2) you have not submitted a Word document. For each visualization question, please make sure that the generated plot is shown in-place with the question and after the code block.

Name and UID

Name: Sib0 Ding

UID: 3035637204

Environmental setup

You need to have the datasets, dplyr, tidyr, rvest, stringr, lubridate, gutenbergr, tidytext, textdata and ggplot2 packages installed. If not yet, please run `install.packages(c("datasets", "tidyr", "dplyr", "rvest", "stringr", "lubridate", "gutenbergr", "tidytext", "textdata", "ggplot2"))` in your R environment.

```
# Load the package.
library(datasets)
library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

1. (3 points) Load the built-in `airquality` dataset and view its first 6 rows. 1) Reshape the dataset (named `airquality_long`) using the `pivot_longer` function to convert the variables `Ozone`, `Solar.R`, `Wind`, and `Temp` into a new column named `Measurement`, with corresponding values in a new column named `Value`. 2) Reshape the `airquality_long` dataset (named `airquality_unite`) using the `unite` function to combine the `Month` and `Day` columns (with `-` as a separator) into a new column named `Date`. Use `head()` to show the results of each sub-question. (hint: you may refer to this link for information: https://www.statology.org/pivot_longer-in-r/)

```
library(tidyr)
```

```
head(airquality) # First 6 rows
```

```
##   Ozone Solar.R Wind Temp Month Day
## 1    41    190  7.4  67    5    1
## 2    36    118  8.0  72    5    2
## 3    12    149 12.6  74    5    3
## 4    18    313 11.5  62    5    4
## 5    NA     NA 14.3  56    5    5
## 6    28     NA 14.9  66    5    6
```

```
# 1)
```

```
airquality_long <- airquality |> pivot_longer(
  cols = c(Ozone, Solar.R, Wind, Temp),
  names_to = "Measurement", values_to = "Value")
head(airquality_long)
```

```
## # A tibble: 6 × 4
##   Month   Day Measurement Value
##   <int> <int>   <chr>      <dbl>
## 1     5     1 Ozone         41
## 2     5     1 Solar.R       190
## 3     5     1 Wind           7.4
## 4     5     1 Temp           67
## 5     5     2 Ozone         36
## 6     5     2 Solar.R       118
```

```
# 2)
airquality_unite <- airquality_long |> unite(col = "Date", c(Month,
Day), sep = "-")
head(airquality_unite)

## # A tibble: 6 × 3
##   Date      Measurement Value
##   <chr>    <chr>         <dbl>
## 1 5-1      Ozone           41
## 2 5-1      Solar.R         190
## 3 5-1      Wind             7.4
## 4 5-1      Temp             67
## 5 5-2      Ozone           36
## 6 5-2      Solar.R         118
```

2. (3 points) Join the following customers and orders data frames by customer_id, with different join function, including: left_join, right_join, inner_join, full_join, semi_join, anti_join (separately), and print the corresponding results (named left_join_df, right_join_df, inner_join_df, full_join_df, semi_join_df and anti_join_df respectively). (hint: <https://www.rdocumentation.org/packages/dplyr/versions/0.7.8/topics/join>, <https://dplyr.tidyverse.org/reference/mutate-joins.html>)

```
customers <- data.frame(
  customer_id = c(1, 2, 3, 4, 5),
  customer_name = c("Alice", "Bob", "Charlie", "Dave", "Eve"),
  city = c("New York", "San Francisco", "Boston", "Seattle", "Chicago")
)
orders <- data.frame(
  customer_id = c(1, 1, 2, 2, 2, 3, 3, 4, 5),
  order_id = c(101, 102, 201, 202, 203, 301, 302, 401, 501),
  order_amount = c(100, 200, 150, 75, 225, 300, 225, 175, 250)
)
```

If there are multiple matches between Left and Right, all combinations of the matches are returned.

```
left_join_df <- left_join(customers, orders, by = "customer_id")
left_join_df
```

```
##   customer_id customer_name      city order_id order_amount
## 1           1         Alice   New York      101          100
## 2           1         Alice   New York      102          200
## 3           2          Bob San Francisco      201          150
## 4           2          Bob San Francisco      202           75
## 5           2          Bob San Francisco      203          225
## 6           3        Charlie    Boston      301          300
## 7           3        Charlie    Boston      302          225
## 8           4          Dave   Seattle      401          175
## 9           5          Eve    Chicago      501          250
```

```
right_join_df <- right_join(customers, orders, by = "customer_id")
right_join_df
```

```
##   customer_id customer_name      city order_id order_amount
## 1           1         Alice   New York     101          100
## 2           1         Alice   New York     102          200
## 3           2           Bob San Francisco     201          150
## 4           2           Bob San Francisco     202           75
## 5           2           Bob San Francisco     203          225
## 6           3         Charlie    Boston     301          300
## 7           3         Charlie    Boston     302          225
## 8           4           Dave   Seattle     401          175
## 9           5           Eve    Chicago     501          250
```

```
inner_join_df <- inner_join(customers, orders, by = "customer_id")
inner_join_df
```

```
##   customer_id customer_name      city order_id order_amount
## 1           1         Alice   New York     101          100
## 2           1         Alice   New York     102          200
## 3           2           Bob San Francisco     201          150
## 4           2           Bob San Francisco     202           75
## 5           2           Bob San Francisco     203          225
## 6           3         Charlie    Boston     301          300
## 7           3         Charlie    Boston     302          225
## 8           4           Dave   Seattle     401          175
## 9           5           Eve    Chicago     501          250
```

```
full_join_df <- full_join(customers, orders, by = "customer_id")
full_join_df
```

```
##   customer_id customer_name      city order_id order_amount
## 1           1         Alice   New York     101          100
## 2           1         Alice   New York     102          200
## 3           2           Bob San Francisco     201          150
## 4           2           Bob San Francisco     202           75
## 5           2           Bob San Francisco     203          225
## 6           3         Charlie    Boston     301          300
## 7           3         Charlie    Boston     302          225
## 8           4           Dave   Seattle     401          175
## 9           5           Eve    Chicago     501          250
```

Return all rows from Left where there are matching values in Right, keeping just columns from Left.

```
semi_join_df <- semi_join(customers, orders, by = "customer_id")
semi_join_df
```

```
##   customer_id customer_name      city
## 1           1         Alice   New York
## 2           2           Bob San Francisco
## 3           3         Charlie    Boston
```

```
## 4          4      Dave      Seattle
## 5          5      Eve       Chicago
```

Return all rows from Left where there are not matching values in Right, keeping just columns from Left.

```
anti_join_df <- anti_join(customers, orders, by = "customer_id")
anti_join_df
```

```
## [1] customer_id  customer_name city
## <0 rows> (or 0-length row.names)
```

3. (2 points) Find the union, intersection and difference of the following df1 and df2 data frames, and print the corresponding results (named union_df, intersect_df, setdiff_df_1_2 and setdiff_df_2_1 respectively).

```
df1 <- data.frame(id = c(1, 2, 3), value = c("a", "b", "c"))
df2 <- data.frame(id = c(3, 4, 5), value = c("c", "d", "e"))
```

```
union_df <- dplyr::union(df1, df2)
union_df
```

```
##   id value
## 1  1     a
## 2  2     b
## 3  3     c
## 4  4     d
## 5  5     e
```

```
intersect_df <- dplyr::intersect(df1, df2)
intersect_df
```

```
##   id value
## 1  3     c
```

```
setdiff_df_1_2 <- dplyr::setdiff(df1, df2)
setdiff_df_1_2
```

```
##   id value
## 1  1     a
## 2  2     b
```

```
setdiff_df_2_1 <- dplyr::setdiff(df2, df1)
setdiff_df_2_1
```

```
##   id value
## 1  4     d
## 2  5     e
```

4. (3 points) Scrape the 1) movie titles, 2) their ratings, and 3) release years from the IMDb Top Rated Movies webpage (<https://www.imdb.com/chart/top/>) with the rvest package. Store the data in a data frame (named movies) and print the top 10 observations in movies. (hint: <https://jtr13.github.io/cc19/web-scraping-using-rvest.html>)

```
library(rvest)
library(stringr)

url <- "https://www.imdb.com/chart/top/"
h <- read_html(url)
tab <- h |> html_nodes("table")
movies <- tab[[1]] |> html_table() # Get the first table

movies <- movies |> select("Rank & Title", "IMDb Rating") |>
  rename(title = "Rank & Title", rating = "IMDb Rating")
# Extract release year within ( )
movies <- movies |> mutate(release_year =
  title |> str_extract("\\d{4}\\") |> str_replace("\\", ""))
# Remove everything before "\n" and after "\n"
movies <- movies |> mutate(title = title |>
  str_replace(".*\\n\\s{6}", "") |> str_replace("\\n.*", ""))
head(movies, 10)

## # A tibble: 10 × 3
##   title                                     rating
##   <chr>                                     <dbl> <chr>
## 1 The Shawshank Redemption                 9.2 1994
## 2 The Godfather                           9.2 1972
## 3 The Dark Knight                         9   2008
## 4 The Godfather Part II                   9   1974
## 5 12 Angry Men                           9   1957
## 6 Schindler's List                       8.9 1993
## 7 The Lord of the Rings: The Return of the King 8.9 2003
## 8 Pulp Fiction                           8.8 1994
## 9 The Lord of the Rings: The Fellowship of the Ring 8.8 2001
## 10 Il buono, il brutto, il cattivo        8.8 1966
```

5. (3 points) Using the stringr package in R, perform the following tasks: 1) Extract all the phone numbers from the following text: "Please call us at 123-456-7890 or 555-555-5555." 2) Extract all the email addresses from the following text: "Contact us at info@example.com or support@example.com." 3) Replace all the URLs (<https://www.xxx.com>) in the following text with the string "URL": "Check out our website at <https://www.example.com> and our blog at <https://blog.example.com>". Print the corresponding results.

```
library(stringr)

# \\d: digit, {}: numbers of occurrence
```

```

"Please call us at 123-456-7890 or 555-555-5555." |>
  str_extract_all("\\d{3}-\\d{3}-\\d{4}")

## [[1]]
## [1] "123-456-7890" "555-555-5555"

# \\w: word character; +: one or more occurrences
"Contact us at info@example.com or support@example.com." |>
  str_extract_all("\\w+@\\w+\\.\\w+")

## [[1]]
## [1] "info@example.com" "support@example.com"

"Check out our website at https://www.example.com and our blog at
https://blog.example.com." |>
  str_replace_all("\\w+://\\w+\\.\\w+\\.\\w+", "URL")

## [1] "Check out our website at URL and our blog at URL."

```

6. (2 points) Using the lubridate package in R, parse the date_time column in the date_data and create new columns for standard date and time components, and print the final results.

```

library(lubridate)

## Loading required package: timechange

##
## Attaching package: 'lubridate'

## The following objects are masked from 'package:base':
##
##   date, intersect, setdiff, union

library(hms)

##
## Attaching package: 'hms'

## The following object is masked from 'package:lubridate':
##
##   hms

date_data <- data.frame(date_time = c("2023-02-22 7:30:15", "2023-02-23
12:15:30", "2023-02-24 23:59:59"))

date_data <- date_data |> mutate(date = date(date_time))
date_data <- date_data |> mutate(time = as_hms(ymd_hms(date_time)))
date_data

##           date_time           date           time
## 1 2023-02-22 7:30:15 2023-02-22 07:30:15

```

```
## 2 2023-02-23 12:15:30 2023-02-23 12:15:30
## 3 2023-02-24 23:59:59 2023-02-24 23:59:59
```

7. (17 points) Explore the advanced data wrangling with the gutenbergr package and its corresponding datasets, and answer the following questions.

a. (1 points) Install the gutenbergr package and load the gutenbergr_metadata as books. Print the first 6 rows, the number of observations (rows) and variables (columns), and the names of all variables in books.

```
library(gutenbergr)

books <- gutenbergr_metadata
head(books) # First 6 rows

## # A tibble: 6 × 8
##   gutenbergr_id title author guten...1 langu...2 guten...3
##   rights has_t...4
##   <chr> <lgl> <int> <chr> <chr> <chr>
## 1 1 "The Declaration o... Jeffe... 1638 en Politi...
##   Publi... TRUE
## 2 2 "The United States... Unite... 1 en Politi...
##   Publi... TRUE
## 3 3 "John F. Kennedy's... Kenne... 1666 en <NA>
##   Publi... TRUE
## 4 4 "Lincoln's Gettysb... Linco... 3 en US Civ...
##   Publi... TRUE
## 5 5 "The United States... Unite... 1 en United...
##   Publi... TRUE
## 6 6 "Give Me Liberty o... Henry... 4 en Americ...
##   Publi... TRUE
## # ... with abbreviated variable names 1gutenberg_author_id, 2language,
## # 3gutenberg_bookshelf, 4has_text

nrow(books) # Number of observations

## [1] 69199

ncol(books) # Number of variables

## [1] 8

names(books) # Names of variables

## [1] "gutenberg_id" "title" "author"
## [4] "gutenberg_author_id" "language"
## [5] "gutenberg_bookshelf"
## [7] "rights" "has_text"
```


b. (2 points) Remove any rows in books that have missing values in the author column, and then count the number of books for each author in a descending order. Who has the most publications and what's the exact number (ignoring Various and Anonymous as an author name)?

```
books |> filter(!is.na(author)) |> # Keep not N/A
  group_by(author) |> count() |>
  arrange(desc(n)) |> head()
```

```
## # A tibble: 6 × 2
## # Groups:   author [6]
##   author                                n
##   <chr>                                <int>
## 1 Various                                3798
## 2 Anonymous                              867
## 3 Shakespeare, William                  326
## 4 Twain, Mark                           235
## 5 Lytton, Edward Bulwer Lytton, Baron   223
## 6 Ebers, Georg                          175
```

Shakespeare, William has the most publications (326).

c. (2 points) Create a subset of books with only Shakespeare, William's English publications, named shakespeare_books. Print the first 6 rows in shakespeare_books.

```
shakespeare_books <- books |> filter(author == "Shakespeare, William" &
  language == "en")
head(shakespeare_books)
```

```
## # A tibble: 6 × 8
##   gutenber...1  title                                author guten...1 langu...2 guten...3
##   rights has_t...4                                <chr>      <int> <chr>      <chr>
##   <chr>  <lg1>
## 1      100 The Complete Works... Shake...      65 en      Plays
##   Publi... TRUE
## 2     1041 Shakespeare's Sonn... Shake...      65 en      <NA>
##   Publi... TRUE
## 3     1045 Venus and Adonis      Shake...      65 en      <NA>
##   Publi... TRUE
## 4     1100 The First Part of ... Shake...      65 en      <NA>
##   Copyr... TRUE
## 5     1101 The Second Part of... Shake...      65 en      <NA>
##   Copyr... TRUE
## 6     1102 The Third Part of ... Shake...      65 en      <NA>
##   Copyr... TRUE
## # ... with abbreviated variable names 1gutenberg_author_id, 2language,
## # 3gutenberg_bookshelf, 4has_text
```

d. (4 points) Filter the dataset `shakespeare_books` to only include specifically the book `Hamlet` as `shakespeare_hamlet`, and extract only `gutenberg_id`, `title` and `author` columns to save, and if there are more than one observation in `shakespeare_hamlet`, just preserve the first observation with `slice()`. Then use `gutenberg_download()` to download the corresponding texts according to `shakespeare_hamlet$gutenberg_id` as `hamlet_text`. Lastly join `shakespeare_hamlet` and `hamlet_text` with `left_join()` as `hamlet_data`, and remove any missing values in the `text` column as well as convert the `text` column to lowercase.

```
shakespeare_hamlet <- shakespeare_books |> filter(title == "Hamlet") |>
  select(gutenberg_id, title, author) |> slice(1)
```

Then

```
hamlet_text <- gutenberg_download(shakespeare_hamlet$gutenberg_id)
```

```
## Determining mirror for Project Gutenberg from
https://www.gutenberg.org/robot/harvest
```

```
## Using mirror http://aleph.gutenberg.org
```

Lastly

```
hamlet_data <- left_join(shakespeare_hamlet, hamlet_text, by =
  "gutenberg_id") |>
  filter(!is.na(text)) |>
  mutate(text = tolower(text))
```

```
head(hamlet_data)
```

```
## # A tibble: 6 × 4
```

```
##   gutenberg_id title   author      text
##   <int> <chr>   <chr>      <chr>
```

```
## 1      1787 Hamlet Shakespeare, William
*****...
```

```
## 2      1787 Hamlet Shakespeare, William this ebook was one of
project gutenb...
```

```
## 3      1787 Hamlet Shakespeare, William time when proofing
methods and tools...
```

```
## 4      1787 Hamlet Shakespeare, William is an improved edition of
this title...
```

```
## 5      1787 Hamlet Shakespeare, William (#100) at
https://www.gutenberg.org/...
```

```
## 6      1787 Hamlet Shakespeare, William
*****...
```

e. (4 points) Perform sentiment analysis on `hamlet_data` using the `tidytext` package. First, get the sentiment lexicon `afinn` through `get_sentiments()` using the `textdata` package and store it in `hamlet_sentiments`. Second, extract each token in `text` column of `hamlet_data` with `unnest_tokens()`. Third, remove the stop words with `anti_join()`. Fourth, join it with `hamlet_sentiments` by `inner_join`. Fifth, count the number of the combination of word and its sentiment value in a descending order by using `count(your_data, word, value, sort=TRUE/FALSE)`, saved as `hamlet_words`. (hint: <http://rafalab.dfci.harvard.edu/dsbook/text-mining.html#sentiment-analysis>)

```
library(tidytext)
library(textdata)

hamlet_sentiments <- get_sentiments(lexicon = "afinn") # First

# unnest_tokens(output column name, input column name)
hamlet_words <- hamlet_data |> unnest_tokens(word, text) |> # Second
  anti_join(stop_words, by = "word") |> # Third
  inner_join(hamlet_sentiments, by = "word") |> # Fourth
  count(word, value, sort = TRUE) # Fifth

head(hamlet_words)

## # A tibble: 6 × 3
##   word    value     n
##   <chr>   <dbl> <int>
## 1 love      3     68
## 2 heaven    2     46
## 3 death   -2     37
## 4 ghost   -1     34
## 5 god       1     33
## 6 dead    -3     31
```

f. (4 points) Following question e, please do operations on a dataset copy of `hamlet_words` as `hamlet_top_words` to obtain the results of the top 1 most common word for each `value` group with `group_by(value)` and `top_n(1, n)`, and reorder the results in a descending order of `n`, then create a bar plot with `geom_col()` of the top 1 most common word for each `value` group in `hamlet_words`. Set an appropriate plot title and axis titles.

```
library(ggplot2)

hamlet_top_words <- hamlet_words |>
  group_by(value) |> top_n(1, n) |>
  arrange(desc(n))

hamlet_top_words |> ggplot(aes(x = word, y = value)) +
  geom_col() +
  ggtitle("Most common words for each sentiment value in Hamlet") +
```

```
xlab("Most common words") +  
ylab("Corresponding sentiment value")
```

